

# aws for you

Instance in an Instant

## The Problem

- Students/researchers in data science have no concrete way to choose AWS instance type.
  - A1, T3, T3a, T2, M5, M5a, M4, C5, C5n, C4
- There is computation time vs cost tradeoff.

## The Solution

- Estimate how long the user's algorithm will take to run on various AWS instances.
- Suggest the best instance choice for runtime or cost.

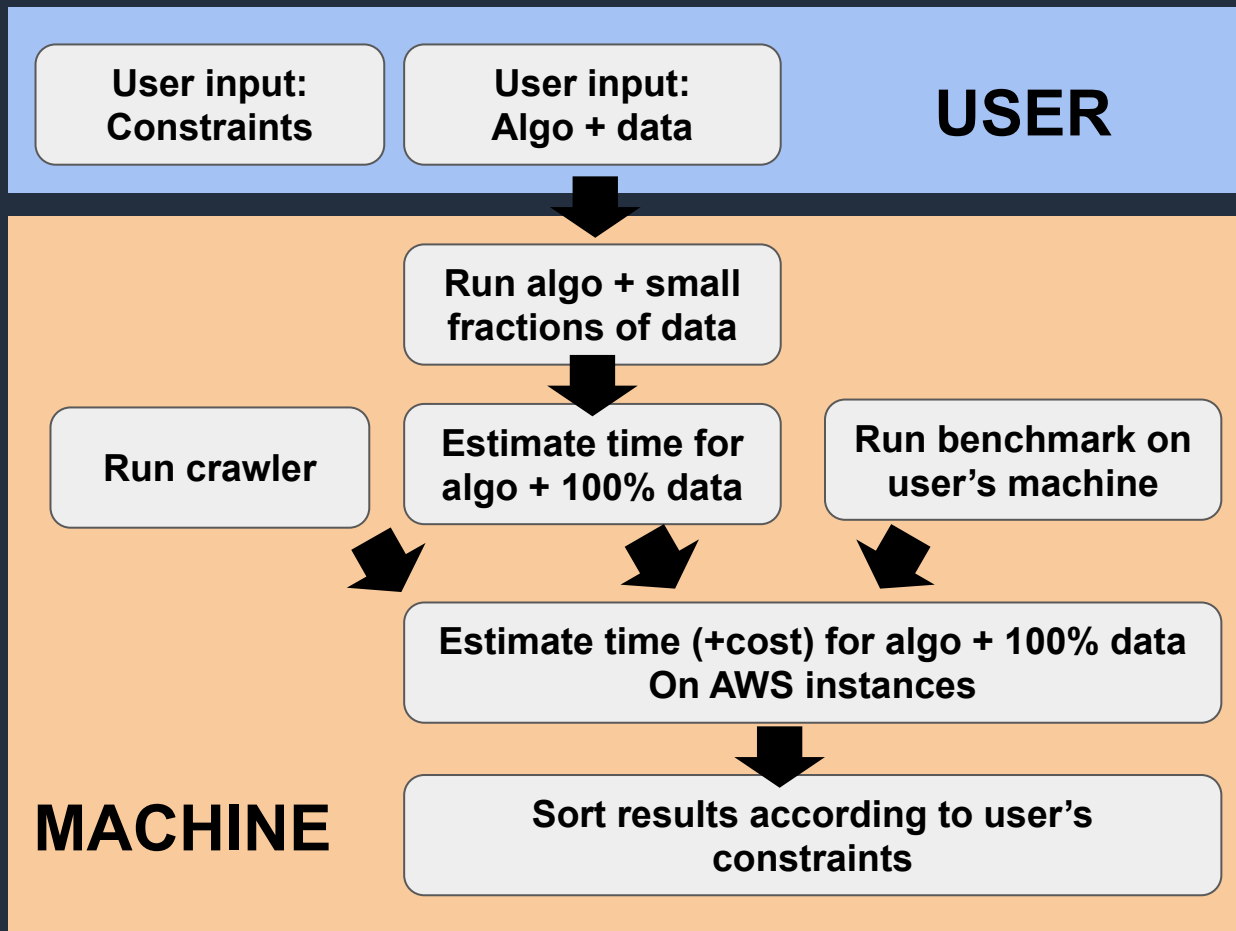
# The Data

- AWS performance data
  - Generated by running benchmark on various types of AWS instances
  - Limitation: New type of AWS instance has to be benchmarked
- On-demand and spot price data
  - Generated by a crawler

```
{  
  'datetime': '28/5/2019 2:06',  
  'RAM': 8,  
  'brand': 'Intel(R) Xeon(R) CPU E5-2686 v4 @ 2.30GHz',  
  'count': 2,  
  'hz_actual': '2.3001 GHz',  
  'hz_advertised': '2.3000 GHz',  
  'instancetype': 't2.large',  
  'region': 'us-east-1',  
  'runtime': 38.95549989,  
}
```

## Use cases

- Choosing the fastest instance given budget.
- Choosing the cheapest instance given time.



# demo

Instance in an Instant

# Components `benchmark_runner`

`command  
line`

`algo_runner`

`total_time`

`recommender`

`report_  
generator`

User input:  
Algo + data

Run algo  
+ small  
fractions  
of data

Estimate  
time for  
algo at  
100% data  
on local  
machine

Estimate time (+cost)  
for algo at 100% data  
On AWS instances

Sort results according  
to user's constraints

User input:  
Constraints

Run benchmark on  
user's machine

Fetch spot prices

`price_crawler`



# algo\_runner

Run algo + small  
fractions of data

- Takes as inputs the python used to call the code in question:

```
"run_mnist(data_loc='data/mnist_data/mnist_data_20k.csv',  
target_loc='data/mnist_data/mnist_target_20k.csv')"
```

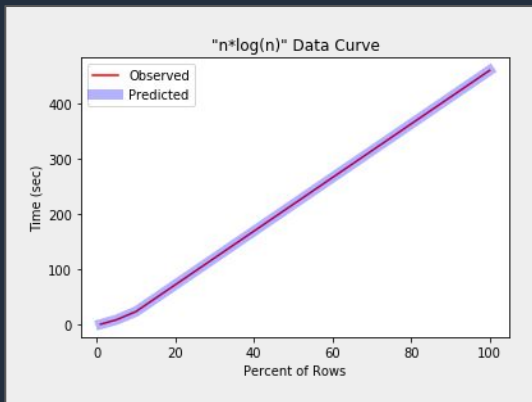
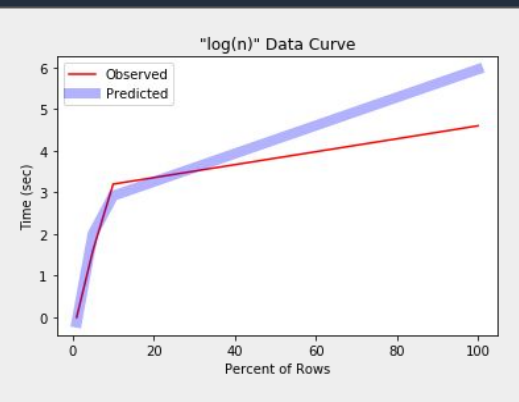
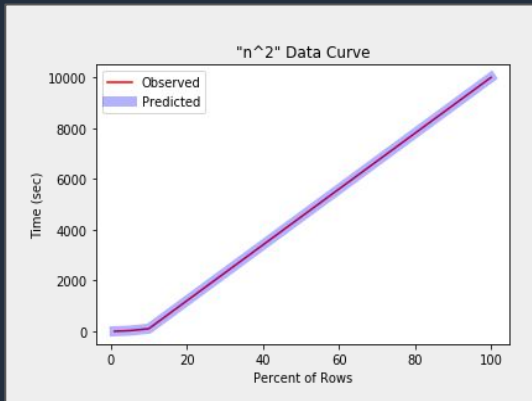

- Outputs two lists, one containing the percentage of samples run per iteration, the other the amount of time (in seconds) it took to run each of the three iterations:

```
([3.1887929439544678, 4.650763988494873, 6.285529851913452], [0.05, 0.1, 0.15])
```

# total\_time

We need to extrapolate the algo runtime at 100% data i.e. estimating complexity.  $O(n, n^2, 1/n, n\log(n))$ .

Estimate time for algo at 100% data on local machine





# benchmark\_runner

Run benchmark on  
user's machine

- Component that runs benchmark on user's machine.
  - Trains a classifier on mnist data (image recognition of digits from 0-9)
  - 60,000 rows for training set, 10,000 rows for test set.
- Also used to run and record the runtime on various AWS instances.

```
>>> import benchmark_runner
Using TensorFlow backend.
>>> benchmark_runner.run_benchmark()
```

```
mnist runtime: 43.468099
43.46809935569763
>>>
```

```
{
  'datetime': '28/5/2019 2:06',
  'RAM': 8,
  'brand': 'Intel(R) Xeon(R) CPU E5-2686 v4 @ 2.30GHz',
  'count': 2,
  'hz_actual': '2.3001 GHz',
  'hz_advertised': '2.3000 GHz',
  'instancetype': 't2.large',
  'region': 'us-east-1',
  'runtime': 38.95549989,
```

# price\_crawler

Fetch spot prices

- Component that fetches the spot price and on-demand price of a given instance at the current instant as a pandas dataframe
  - Get spot price using boto3
  - Get on-demand price using pre aggregated database

	region	spot_price	on_demand_price
0	eu-north-1	0.0130	0.0432
1	ap-south-1	0.0134	0.0448
2	eu-west-3	0.0472	0.0472
3	eu-west-2	0.0142	0.0472
4	eu-west-1	0.0137	0.0456
5	ap-northeast-2	0.0520	0.0520
6	ap-northeast-1	0.0163	0.0544
7	sa-east-1	0.0202	0.0672
8	ca-central-1	0.0139	0.0464
9	ap-southeast-1	0.0158	0.0528
10	ap-southeast-2	0.0158	0.0528
11	eu-central-1	0.0144	0.0480
12	us-east-1	0.0125	0.0416
13	us-east-2	0.0125	0.0416
14	us-west-1	0.0149	0.0496
15	us-west-2	0.0125	0.0416

# recommender

Estimate time (+cost)  
for algo at 100% data  
On AWS instances

Sort results  
according to  
user's constraints

Calls:

- algo\_runner
- algo\_analyzer
- benchmark\_runner
- price\_crawler

Outputs a dataframe to the  
“report\_generator” component, that  
has total estimated times, total  
estimated costs for each instance type.

instance_type	runtime	estimated_time_aws	region	spot_price	on_demand_price	est_cost_spot_price	est_cost_on_demand_price
c5.18xlarge	12.2972	8.5383	eu-north-1	0.9828	3.2760	0.0023	0.0078
c5.18xlarge	12.2972	8.5383	ap-south-1	1.0432	3.0600	0.0025	0.0073
c5.18xlarge	12.2972	8.5383	eu-west-3	1.0908	3.6360	0.0026	0.0086
c5.18xlarge	12.2972	8.5383	eu-west-2	1.1311	3.6360	0.0027	0.0086
c5.18xlarge	12.2972	8.5383	eu-west-1	1.2367	3.4560	0.0029	0.0082
c5.18xlarge	12.2972	8.5383	ap-northeast-2	1.0788	3.4560	0.0026	0.0082
c5.18xlarge	12.2972	8.5383	ap-northeast-1	1.1976	3.8520	0.0028	0.0091
c5.18xlarge	12.2972	8.5383	sa-east-1	1.4685	4.7160	0.0035	0.0112
c5.18xlarge	12.2972	8.5383	ca-central-1	1.0408	3.3480	0.0025	0.0079
c5.18xlarge	12.2972	8.5383	ap-southeast-1	1.1212	3.5280	0.0027	0.0084
c5.18xlarge	12.2972	8.5383	ap-southeast-2	1.2404	3.9960	0.0029	0.0095
c5.18xlarge	12.2972	8.5383	eu-central-1	1.1936	3.4920	0.0028	0.0083
c5.18xlarge	12.2972	8.5383	us-east-1	1.1920	3.0600	0.0028	0.0073
c5.18xlarge	12.2972	8.5383	us-east-2	0.9508	3.0600	0.0023	0.0073
c5.18xlarge	12.2972	8.5383	us-west-1	1.1078	3.8160	0.0026	0.0091
c5.18xlarge	12.2972	8.5383	us-west-2	1.1915	3.0600	0.0028	0.0073
c5.2xlarge	21.8342	15.1600	eu-north-1	0.1092	0.3640	0.0005	0.0015
c5.2xlarge	21.8342	15.1600	ap-south-1	0.1209	0.3400	0.0005	0.0014
c5.2xlarge	21.8342	15.1600	eu-west-3	0.1235	0.4040	0.0005	0.0017
c5.2xlarge	21.8342	15.1600	eu-west-2	0.1257	0.4040	0.0005	0.0017
c5.2xlarge	21.8342	15.1600	eu-west-1	0.1545	0.3840	0.0007	0.0016
c5.2xlarge	21.8342	15.1600	ap-northeast-2	0.1218	0.3840	0.0005	0.0016
c5.2xlarge	21.8342	15.1600	ap-northeast-1	0.1384	0.4280	0.0006	0.0018
c5.2xlarge	21.8342	15.1600	sa-east-1	0.1646	0.5240	0.0007	0.0022
c5.2xlarge	21.8342	15.1600	ca-central-1	0.1156	0.3720	0.0005	0.0016
c5.2xlarge	21.8342	15.1600	ap-southeast-1	0.1276	0.3920	0.0005	0.0017
c5.2xlarge	21.8342	15.1600	ap-southeast-2	0.1389	0.4440	0.0006	0.0019
c5.2xlarge	21.8342	15.1600	eu-central-1	0.1387	0.3880	0.0006	0.0016
c5.2xlarge	21.8342	15.1600	us-east-1	0.1297	0.3400	0.0005	0.0014
c5.2xlarge	21.8342	15.1600	us-east-2	0.0760	0.3400	0.0003	0.0014
c5.2xlarge	21.8342	15.1600	us-west-1	0.1384	0.4240	0.0006	0.0018
c5.2xlarge	21.8342	15.1600	us-west-2	0.1484	0.3400	0.0006	0.0014

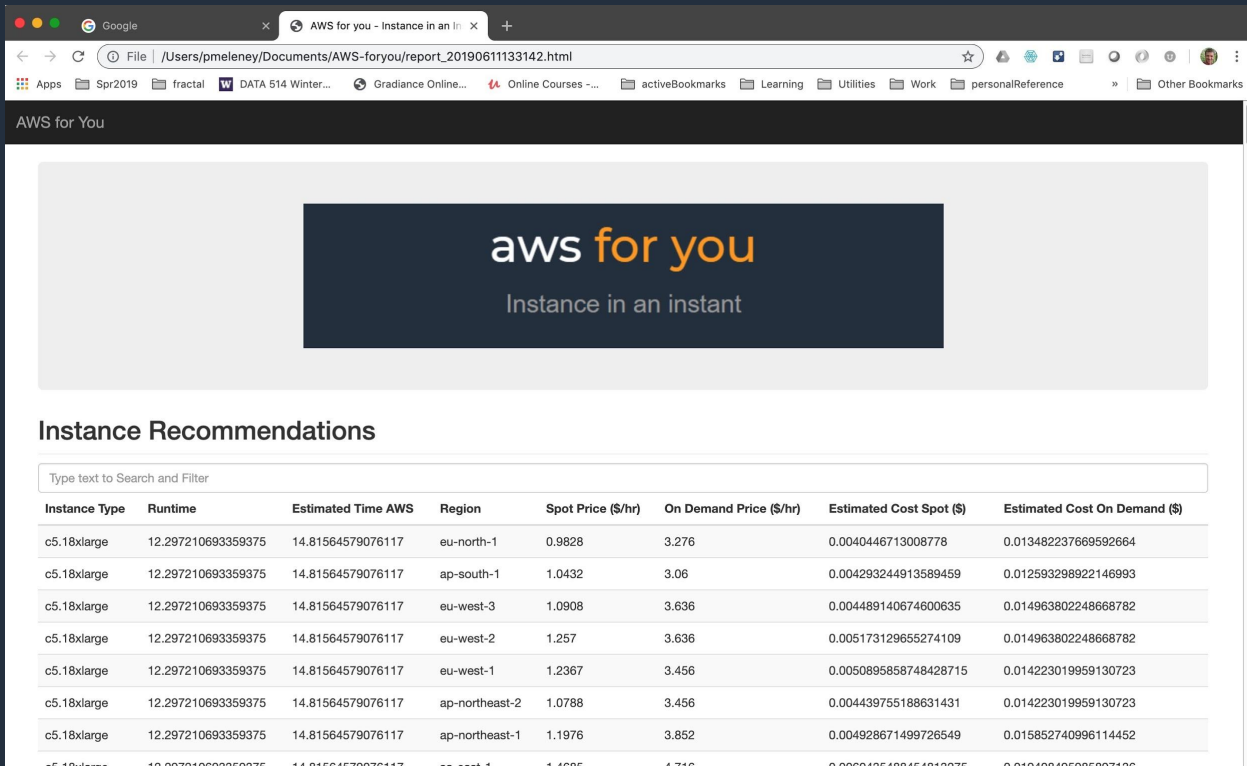
# report\_generator

Takes the dataframe  
output from the  
recommender and  
generates a sortable  
html table

User finds best  
instance for price or  
time by sorting

User input:  
Constraints

User input:  
Algo + data



aws for you

Instance in an instant

### Instance Recommendations

Type text to Search and Filter

Instance Type	Runtime	Estimated Time AWS	Region	Spot Price (\$/hr)	On Demand Price (\$/hr)	Estimated Cost Spot (\$)	Estimated Cost On Demand (\$)
c5.18xlarge	12.297210693359375	14.81564579076117	eu-north-1	0.9828	3.276	0.0040446713008778	0.013482237669592664
c5.18xlarge	12.297210693359375	14.81564579076117	ap-south-1	1.0432	3.06	0.004293244913589459	0.012593289922146993
c5.18xlarge	12.297210693359375	14.81564579076117	eu-west-3	1.0908	3.636	0.004489140674600635	0.014963802248668782
c5.18xlarge	12.297210693359375	14.81564579076117	eu-west-2	1.257	3.636	0.005173129655274109	0.014963802248668782
c5.18xlarge	12.297210693359375	14.81564579076117	eu-west-1	1.2367	3.456	0.0050895858748428715	0.014223019959130723
c5.18xlarge	12.297210693359375	14.81564579076117	ap-northeast-2	1.0788	3.456	0.004439755188631431	0.014223019959130723
c5.18xlarge	12.297210693359375	14.81564579076117	ap-northeast-1	1.1976	3.852	0.004928671499726549	0.015852740996114452
c5.18xlarge	12.297210693359375	14.81564579076117	sa-east-1	1.4685	4.716	0.0060435488454813275	0.019408405085807126

# demo

Instance in an Instant

# Project Structure and Continuous Integration

build **passing**

coverage **99%**

license **MIT**

contributors **4**

code size **31.2 kB**

pull requests **0 open**

closed pull requests **23**

```
AWS-foryou/  
| - README.md  
| - awsforyou/  
|   | - __init__.py  
|   | - algo_runner.py  
|   | - aws_metadata.py  
|   | - aws_pricing.py  
|   | - benchmark_runner.py  
|   | - recommender.py  
|   | - report_generator.py  
|   | - total_time_component.py  
|   | - ui/  
|     | - template.html  
| - tests/  
|   | - __init__.py  
|   | - test_algo_runner.py  
|   | - test_aws_metadata.py  
|   | - test_aws_pricing.py  
|   | - test_benchmark_runner.py  
|   | - test_keras_mnist.py  
|   | - test_recommender.py  
|   | - test_report_generator.py  
|   | - test_total_time_component.py
```

Ran 26 tests in 806.644s

OK

The command "coverage run -m unittest discover awsforyou" exited with 0.

```
| - data/  
|   | - aws-scorecard.csv  
| - docs/  
|   | - component-specification.md  
|   | - functional-specification.md  
| - examples/  
|   | - demo.py  
|   | - examples.ipynb  
|   | - sklearn_diabetes.py  
|   | - x_diabetes.csv  
|   | - y_diabetes.csv  
| - setup.py  
| - requirements.txt  
| - LICENSE
```

# Challenges

- Configuring security credentials in CI
- Correct model selection for time estimate
- AWS python SDK (boto3) API
- Writing tests for code meant to be run on AWS to fetch instance-related metadata (mocking).
- pip installable package

# Future Work

- Ability to recommend GPU instances
- Extend coverage to Azure/Google Cloud



# aws for you

Instance in an instant